

Outline

- Sequential Data
- Intro to Recurrent Neural Networks (RNNs)
- Challenges with RNNs
- Example: Sine Wave Generation
- Assignment: Cosine Waves


Sequential Data

Example Sequential Data and Tasks

- Written Language
 - Character Prediction
 - Machine Translation
- Audio
 - Speech Processing
 - Music Transcription
- Spatio-Temporal Data
 - Body capture data
 - Weather modelling
 - Neurological Models




Language Modelling/Prediction




Machine Translation

Example Sequential Data and Tasks

- Written Language
 - Character Prediction
 - Machine Translation
- Audio
 - Speech Recognition
 - Music Transcription
- Spatio-Temporal Data
 - Body capture data
 - Weather modelling
 - Neurological Models




"Hey Google, what is the weather today?"
Speech Recognition



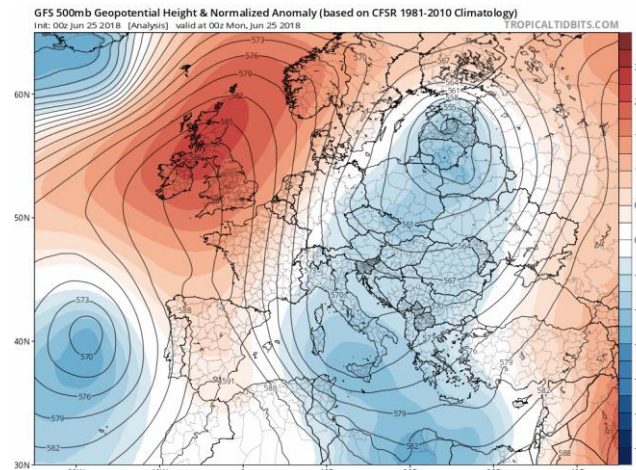
Music Transcription

Example Sequential Data and Tasks

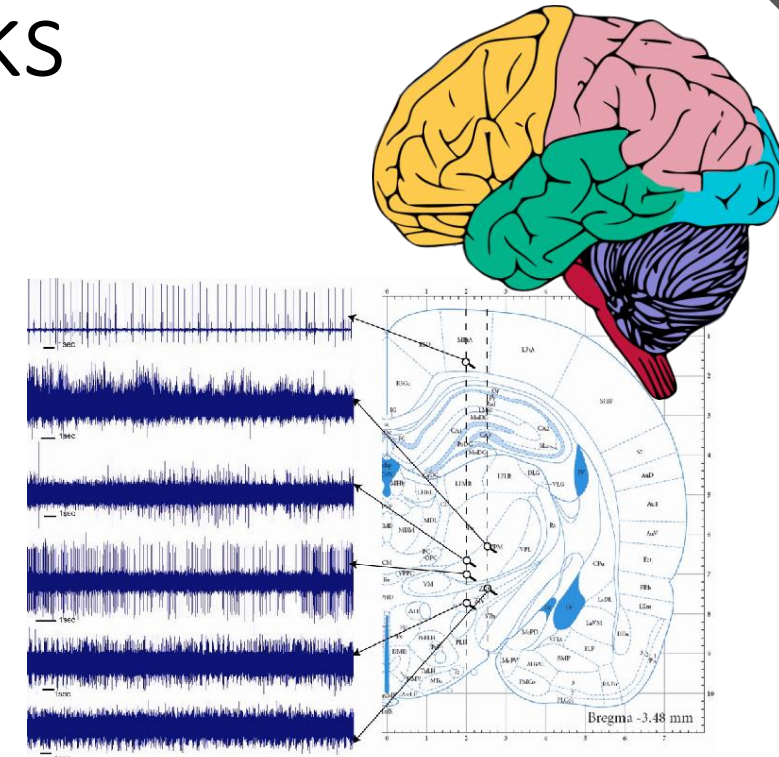
- Written Language
 - Character Prediction
 - Machine Translation
- Audio
 - Speech Recognition
 - Music Transcription
- Spatio-Temporal Data
 - Body capture data
 - Weather modelling
 - Neurological Data



Body Motion Capture



Weather Modelling



Brain Activity

Features of Sequential Data


- Order matters
 - Unlike a group of vectors, the order in which data appear in sequential data is important
- Variable length
 - Measurements are not always captured over the same number of time-steps
- Temporal dependence
 - Previous data values usually has impact on current value



Recurrent Neural Networks

RNN Setup

- Processes input (x) at each step using shared parameters (A)
- Retains memory by evolving the hidden state (h) as a function of input



RNN Setup

- Output of a cell is the hidden state $h^{(t)}$
- Function of input (x) at time t and previous hidden state


$$\mathbf{a}^{(t)} = \mathbf{b} + \mathbf{W}\mathbf{h}^{(t-1)} + \mathbf{U}\mathbf{x}^{(t)}$$

$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)})$$

RNN Implementation

`torch.nn.RNN()`

- Parameters
 - `input_size`
 - `hidden_size`
 - `num_layers`
 - `nonlinearity`
 - `bias`
 - `batch_first`
 - `dropout`
 - `bidirectional`



RNN Implementation

`torch.nn.RNN()`

- Parameters

- `input_size`
- `hidden_size`
- `num_layers`
- `nonlinearity`
- `bias`
- `batch_first`
- `dropout`
- `bidirectional`

- Input size is the dimension of the input at each time step.
 - The size of each vector in a sequence
- Hidden size is the dimensionality of the hidden state, h .
 - The number of neurons processing each time step

RNN Implementation

`torch.nn.RNN()`

- Parameters

- `input_size`
- `hidden_size`
- `num_layers`
- `nonlinearity`
- `bias`
- `batch_first`
- `dropout`
- `bidirectional`

- Num_layers is the number of stacked RNN layers (default: 1)
 - Will be covered in detail in Lab 6
- Non-linearity is the activation function for the neurons
 - Choice of `'tanh'` or `'relu'` (default: `tanh`)

RNN Implementation

`torch.nn.RNN()`

- Parameters

- `input_size`
- `hidden_size`
- `num_layers`
- `nonlinearity`
- `bias`
- `batch_first`
- `dropout`
- `bidirectional`

- Bias is a Boolean indicating whether to include the bias term (b) in the neuron (default: True)

$$a^{(t)} = b + Wh^{(t-1)} + Ux^{(t)}$$

- `Batch_first` indicates whether the first dimension of the input is the batch- (*batch, seq, feature*)
 - Default: False – (*seq, batch, feature*)

RNN Implementation

`torch.nn.RNN()`

- Parameters

- `input_size`
- `hidden_size`
- `num_layers`
- `nonlinearity`
- `bias`
- `batch_first`
- `dropout`
- `bidirectional`

- Dropout: if non-zero, introduces a *Dropout* layer on the outputs of each RNN layer except the last
 - No effect if `num_layers = 1`
- Bidirectional: Boolean indicating whether the RNN is bidirectional (Default: False)

RNN Usage

- Inputs:
 - `input` is a tensor containing the features of the input sequence
 - Default shape: *(seq_len, batch, input_size)* – can change if `batch_first`
 - Can be packed variable length sequence (see [`torch.nn.utils.rnn.pack_padded_sequence\(\)`](#) or [`torch.nn.utils.rnn.pack_sequence\(\)`](#))

RNN Usage

- Inputs:
 - `input` is a tensor containing the features of the input sequence
 - Default shape: $(seq_len, batch, input_size)$ – can change if `batch_first`
 - Can be packed variable length sequence (see [torch.nn.utils.rnn.pack_padded_sequence\(\)](#) or [torch.nn.utils.rnn.pack_sequence\(\)](#))
 - `h_0` is a tensor containing the initial hidden state for each element in batch
 - Shape: $(num_layers * num_directions, batch, hidden_size)$
 - Common practice is to use `torch.zeros(num_layers * num_directions, batch, hidden_size)`
 - If None, will default to `torch.zeros` of the appropriate size

RNN Usage

- Outputs
 - output is a tensor containing the output features (h_t) from the **last** layer of the RNN, for each t .
 - Default shape: $(seq_len, batch, num_directions * hidden_size)$

RNN Usage

- Outputs

- `output` is a tensor containing the output features (h_t) from the last layer of the RNN, for each t .


- Default shape: $(seq_len, batch, num_directions * hidden_size)$

- `h_n` is a tensor containing the all the hidden states for $t = seq_len$ (i.e., after the final time step of the sequence)


- Shape: $(num_layers * num_directions, batch, hidden_size)$

RNN Problem Types


one to one




one to many




many to one




many to many






many to many




RNN Problem Types






RNN Problem Types



RNN Problem Types



RNN Problem Types




RNN Challenges

Vanishing and Exploding Gradient

- To learn on sequential data, we need:
 - Sensitivity to new input
 - Retention of old information
- Backpropagation over large number of steps leads to issues
 - Amplification of patterns by repeated application of same function

Vanishing and Exploding Gradient

- Long-term memory of RNNs suffers due to uncontrolled gradients
- Additional mechanisms are used to optimize memory and sensitivity (details in Lab 6)
 - Gated architectures – LSTM, GRU
 - Constrained-weight RNNs



Sensitivity to Initialization

- RNN performance can be very sensitive to initialization
- Optimal initialization distribution depends on:
 - Hidden size
 - Non-linearity
 - Input statistics

Initialization Distributions in PyTorch

`torch.nn.init`


- Uniform
- Normal
- Xavier Uniform
- Xavier Normal
- Kaiming Uniform
- Kaiming Normal
- Orthogonal

Example: Sine Wave Generation

Adapted this tutorial: <https://lirnli.wordpress.com/2017/09/01/simple-pytorch-rnn-examples/>

Problem setup

- Sine wave with random shift as input
- Want to predict the next time step given an input sequence




Model Definition

- Using RNN with hidden size 128
- Use linear layer for output
- Teacher forcing to help guide network with true values
 - hybrid one-to-many/many-to-many problem

```
1 # Model Definition
2 class SineRNN(nn.Module):
3     def __init__(self, p = 0.5):
4         super(SineRNN, self).__init__()
5         self.rnn_layer = nn.RNN(input_size = 1, hidden_size = 128)
6         self.out_layer = nn.Linear(in_features = 128, out_features = 1)
7         self.p = p #Whether to use actual seq or output for next step
8
9     def forward(self, seq, h = None):
10         out = []
11         X_in = torch.unsqueeze(seq[0], 0)
12         for X in seq:
13             if np.random.rand() > self.p: #Use teacher forcing
14                 X_in = X.unsqueeze(dim = 0)
15             tmp, h1 = self.rnn_layer(X_in, h1)
16             X_in = self.out_layer(tmp)
17             out.append(X_in)
18         return torch.stack(out).squeeze(1), h1
```

Teacher Forcing

- When predicting a sequence from an input sequence, you can use “teacher forcing”
- Randomly select either:
 - Actual input (teacher) value
 - Output of network at previous step (prediction)



Teacher Forcing


Training

- Generate data:
 - Sine waves with random shifts
 - Both input and targets (shifted 1)
- Increase chance of using past predictions
- Warm-up for 20 steps and then evaluate loss

```
1 seq = SineRNN()
2 criterion = nn.MSELoss()
3 optimizer = optim.Adam(seq.parameters(), lr=0.001)
4 max_iters = 10000
5 train_loss = []
6 for i in range(max_iters):
7     data = np.sin(np.linspace(0,10,100)+2*np.pi*np.random.rand())
8     xs = data[:-1]
9     ys = data[1:]
10    X = torch.Tensor(xs).view(-1,1,1)
11    y = torch.Tensor(ys)
12    if i%100==0:
13        seq.p = min(seq.p+0.1,0.85) # encourage training longer term predictions
14    optimizer.zero_grad()
15    rnn_out,_ = seq(X)
16    loss = criterion(rnn_out[20:].view(-1),y[20:])
17    loss.backward()
18    optimizer.step()
19    train_loss.append(loss.item())
20    if i%500 == 0:
21        print(f"i {i}, loss {loss.data:.4f}")
```

Results- Predictions

- First 100 steps as seed
- Predict for 1000 steps
- Able to capture general shape, but makes errors finding period (requires long memory)

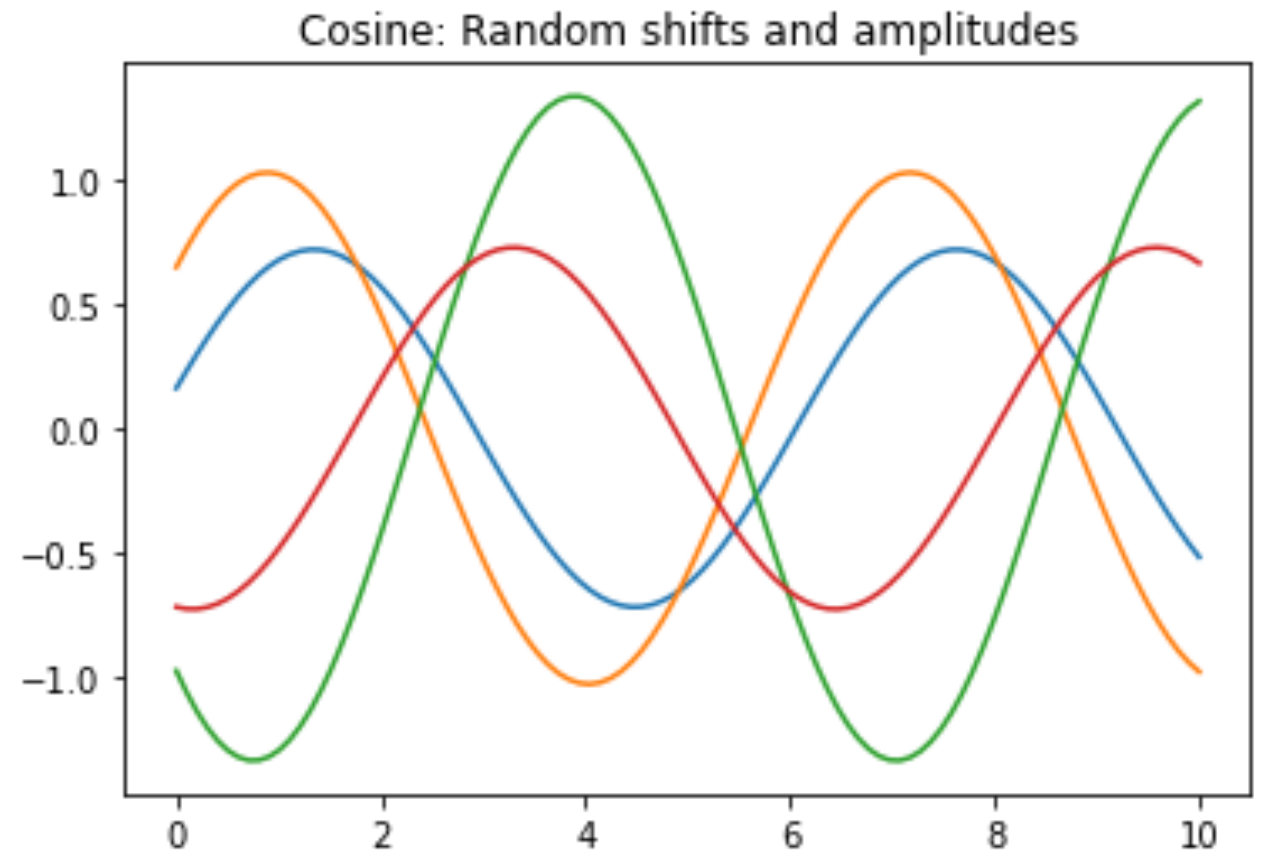


Total MSE Loss: 0.9695

Assignment: Cosine Wave Generation


Assignment Details

- Create an RNN model to generate cosines like in example
- Use random shifts as in example
- Generate samples with different amplitudes, ranging from 0.5 to 1.5 (How much data do you need?)
- Play with hyperparameters of your network



Evaluation

- Generate a validation/test set of 1000 randomly-generated cosines
- Find MSE loss on this validation set
- Plot your best and worst predictions in the validation set



Evaluation Detail

- For validation set, create longer sequences (1000 steps)
- Evaluate loss after a warm-up period of 100 steps
- Only calculate loss where prediction and real input overlap

