

# Lab 6: More on RNNs

University of Washington

ECE 596/AMATH 563

Spring 2021

# Outline

## Part 1: Gated RNN Architectures

- LSTM
- GRU

## Part 2: Examples of Different RNN Problems

- One-to-one
- One-to-many
- Many-to-one
- Many-to-many


## Part 3: Other RNN Variants

- Deep RNN
- Bidirectional RNN

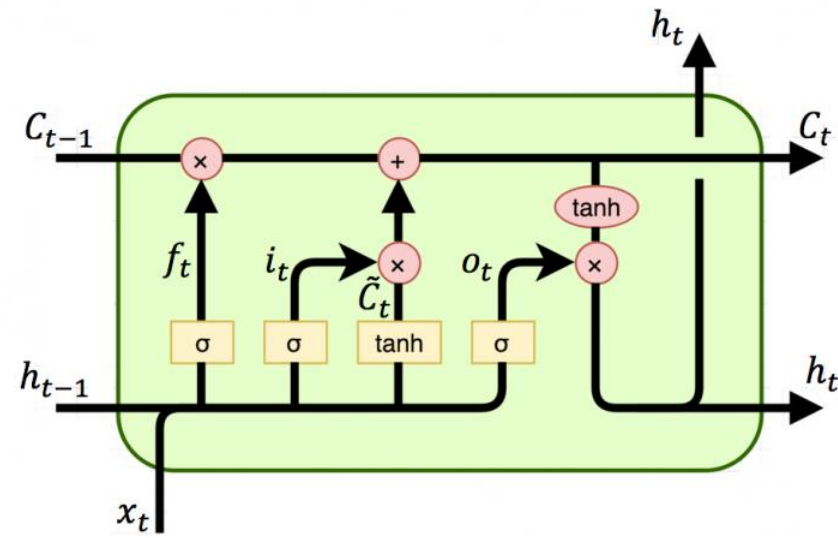
## Lab Assignment

# Part 1: Gated RNN Architectures


# LSTM (Long Short-Term Memory)




**Vanilla RNN**




**LSTM**

  
Neural Network  
Layer


  
Pointwise  
Operation

  
Vector  
Transfer

  
Concatenate

  
Copy


# LSTM: Detailed Architecture



## Cell state


- Unique to LSTM
- Long term memory of the model

# LSTM: Detailed Architecture



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

**Forget gate layer**

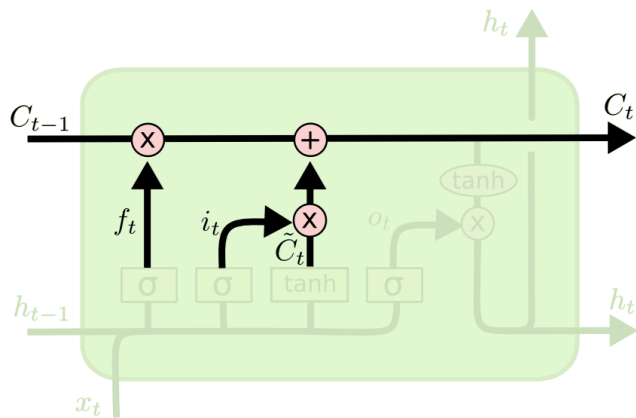


$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$


**Input gate layer**

# LSTM: Detailed Architecture



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

**Update cell state**



$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

**Output gate layer**

# LSTM: PyTorch Implementation

```
class LSTMModel(nn.Module):
    def __init__(self, input_dim, hidden_dim, layer_dim, output_dim):
        super(LSTMModel, self).__init__()

        self.hidden_dim = hidden_dim
        self.layer_dim = layer_dim
        self.lstm = nn.LSTM(input_dim, hidden_dim, layer_dim)
        self.fc = nn.Linear(hidden_dim, output_dim)

    def forward(self, x):
        h0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim).requires_grad_()
        c0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim).requires_grad_()
        out, (hn, cn) = self.lstm(x, (h0, c0))
        out = self.fc(out[:, -1, :])
        return out
```

**Hidden layer dimension**

**# of hidden layers**

**# PyTorch built in LSTM architecture**

**Readout layer**


**Initialize hidden/cell states**

**Pass the input x to LSTM**  
(x.shape = seq\_len, batch, input\_size)


**Pass the hidden state of last time step to readout layer**






# GRU (Gated Recurrent Units)




**LSTM**



**GRU**



# GRU: Detailed Architecture



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# GRU: PyTorch Implementation

```
class GRUModel(nn.Module):
    def __init__(self, input_dim, hidden_dim, num_layers, output_dim):
        super(LSTMModel, self).__init__()

        self.hidden_dim = hidden_dim
        self.num_layers = num_layers

        self.gru = nn.GRU(input_dim, hidden_dim, num_layers)
        self.fc = nn.Linear(hidden_dim, output_dim)

    def forward(self, x):
        h0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim).requires_grad_()

        out, (hn, cn) = self.gru(x, h0)
        out = self.fc(out[:, -1, :])

        return out
```

**Hidden layer dimension**

**# of hidden layers**

**# PyTorch built in GRU architecture**

**Readout layer**

**Pass the input x to GRU**


**Initialize hidden/cell states**

**Pass the hidden state of last time step to readout layer**


# Part 2: Examples of Different RNN Problems

# One-to-one


one to one




one to many




many to one



many to many




many to many




# One-to-Many


one to one




one to many




many to one



many to many



many to many



# One-to-Many


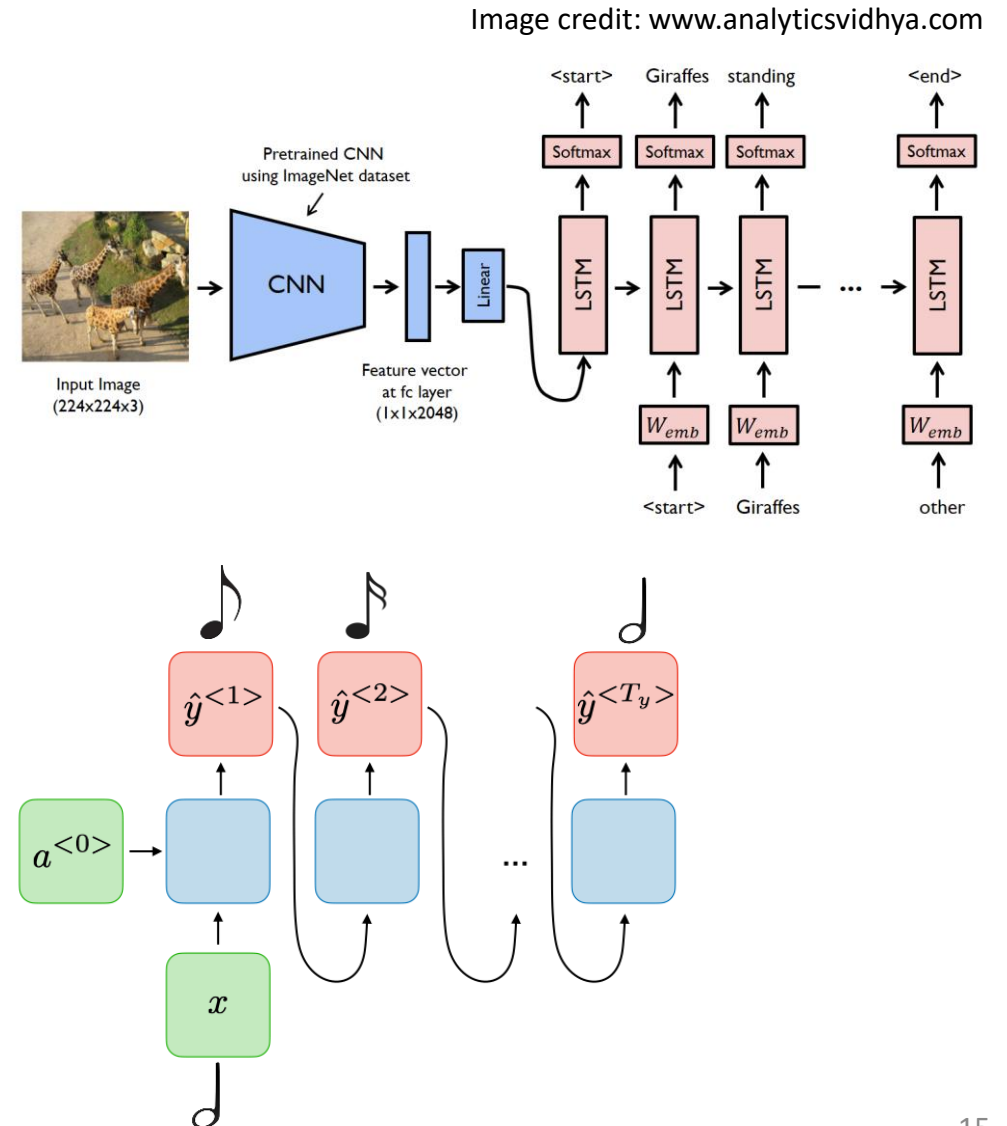


Image captioning

Music generation



# One-to-Many: Image captioning example

```
class DecoderRNN(nn.Module):
    def __init__(self, embed_size, hidden_size, vocab_size, num_layers=1):
        super(DecoderRNN, self).__init__()


        # define the properties
        self.embed_size = embed_size
        self.hidden_size = hidden_size
        self.vocab_size = vocab_size

        # lstm cell
        self.lstm_cell = nn.LSTMCell(input_size=embed_size, hidden_size=hidden_size)

        # output fully connected layer
        self.fc_out = nn.Linear(in_features=self.hidden_size, out_features=self.vocab_size)

        # embedding layer
        self.embed = nn.Embedding(num_embeddings=self.vocab_size, embedding_dim=self.embed_size)

        # activations
        self.softmax = nn.Softmax(dim=1)
```



## Word Embedding


Word	Integer	Vector representation
apple	1	[0 0 0 0 1]
orange	2	[0 0 0 1 0]
guava	3	[0 0 0 1 1]
blue	4	[0 0 1 0 0]
green	5	[0 0 1 0 1]
red	6	[0 0 1 1 0]

Image credit: Medium



# One-to-Many: Image captioning example

```
def forward(self, features, captions):  
    # batch size  
    batch_size = features.size(0)  
  
    # init the hidden and cell states to zeros  
    hidden_state = torch.zeros((batch_size, self.hidden_size)).cuda()  
    cell_state = torch.zeros((batch_size, self.hidden_size)).cuda()  
  
    # define the output tensor placeholder  
    outputs = torch.empty((batch_size, captions.size(1), self.vocab_size)).cuda()  
  
    # embed the captions  
    captions_embed = self.embed(captions)  
  
    # pass the caption word by word  
    for t in range(captions.size(1)):  
        # for the first time step the input is the feature vector  
        if t == 0:  
            hidden_state, cell_state = self.lstm_cell(features, (hidden_state, cell_state))  
  
        # for the 2nd+ time step, using teacher forcer  
        else:  
            hidden_state, cell_state = self.lstm_cell(captions_embed[:, t, :], (hidden_state, cell_state))  
  
        # output of the attention mechanism  
        out = self.fc_out(hidden_state)  
  
        # build the output tensor  
        outputs[:, t, :] = out  
  
    return outputs
```




Implementation detail:


<https://medium.com/@stepanulyanin/captioning-images-with-pytorch-bc592e5fd1a3>

# Many-to-One


one to one




one to many




many to one



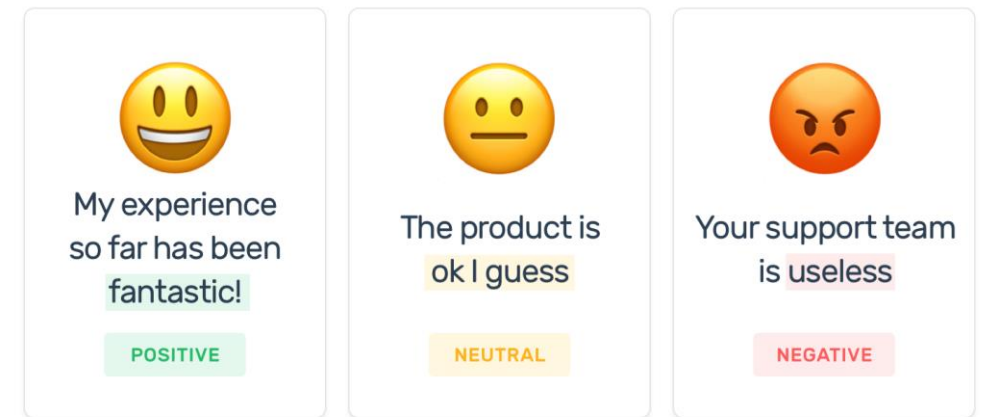
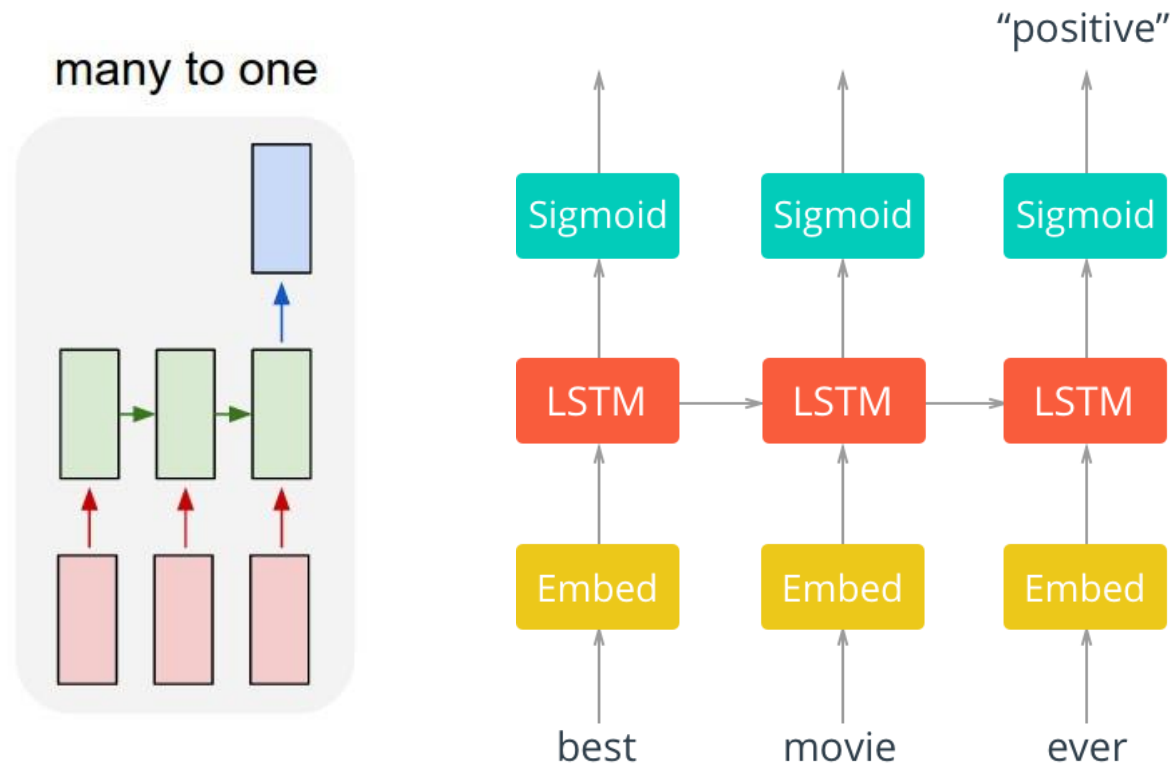
many to many



many to many



# Many-to-One



## Sentiment Analysis

# Many-to-one: Simple sentiment analysis

```
import torch.nn as nn

class RNN(nn.Module):
    def __init__(self, input_dim, embedding_dim, hidden_dim, output_dim):
        super().__init__()

        self.embedding = nn.Embedding(input_dim, embedding_dim)

        self.rnn = nn.RNN(embedding_dim, hidden_dim)

        self.fc = nn.Linear(hidden_dim, output_dim)

    def forward(self, text):
        #text = [sent len, batch size]

        embedded = self.embedding(text)

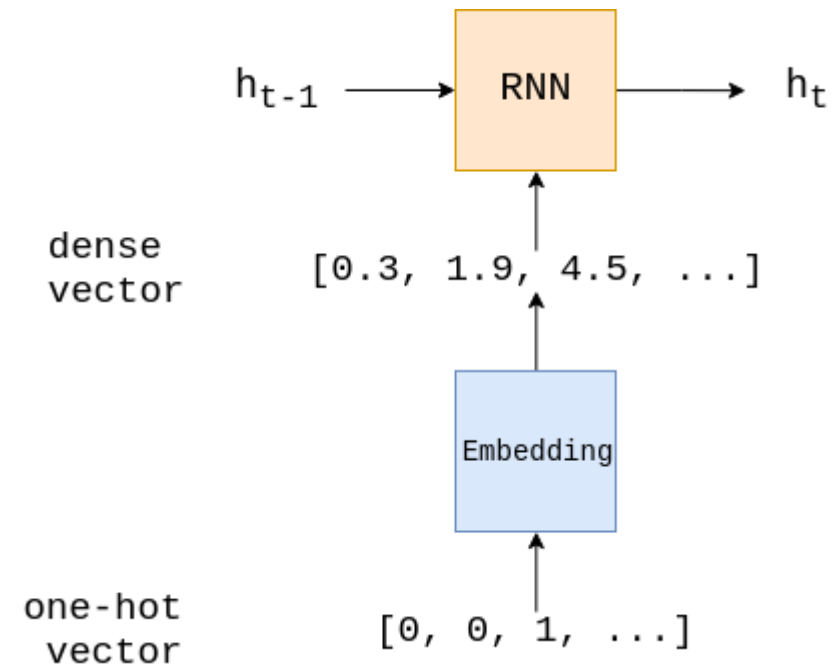
        #embedded = [sent len, batch size, emb dim]

        output, hidden = self.rnn(embedded)

        #output = [sent len, batch size, hid dim]
        #hidden = [1, batch size, hid dim]

        assert torch.equal(output[-1, :, :], hidden.squeeze(0))


        return self.fc(hidden.squeeze(0))
```




Implementation detail :  
<https://towardsdatascience.com/sentiment-analysis-using-lstm-step-by-step-50d074f09948>

# Many-to-Many


one to one




one to many




many to one



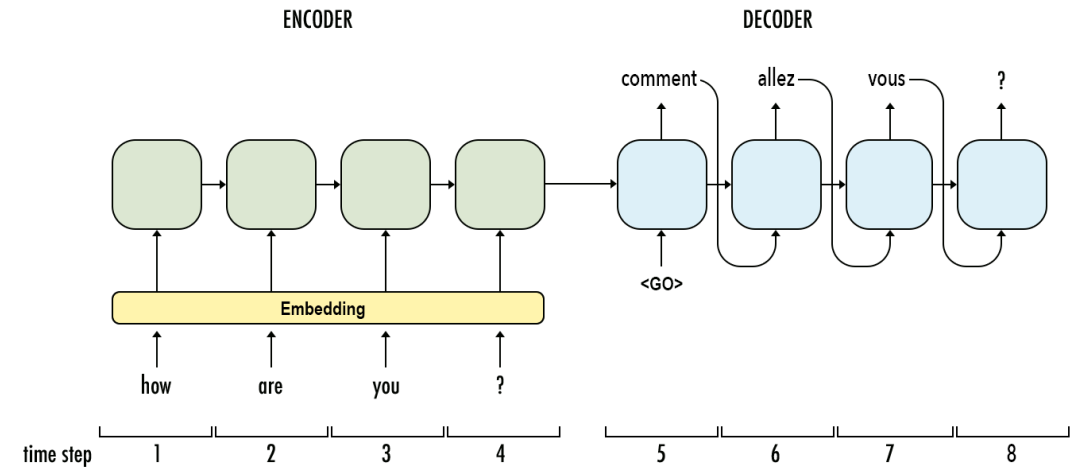

many to many




many to many



# Many-to-Many



## Machine Translation



## Video Captioning

# Many-to-Many: Machine translation example

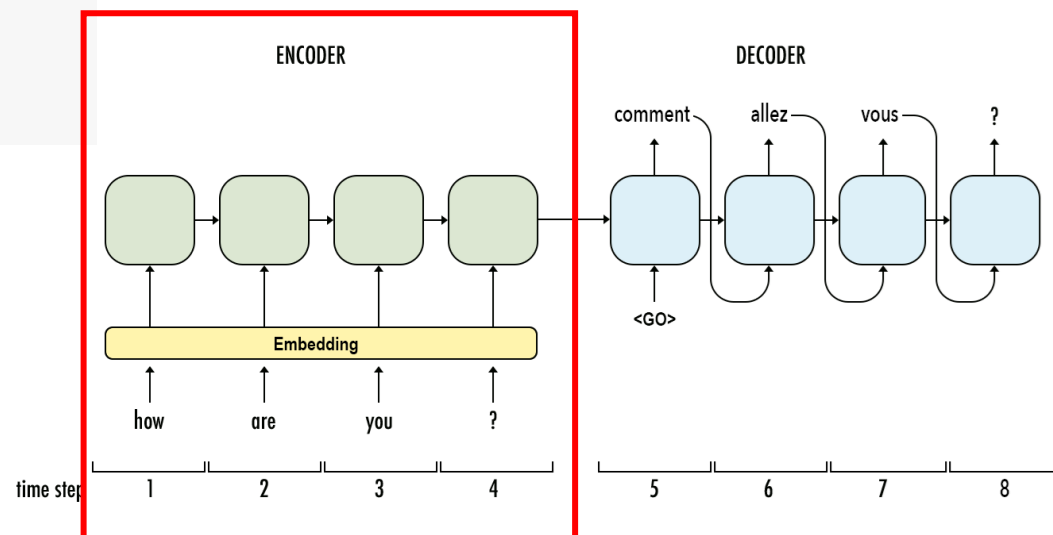
```
class Encoder(nn.Module):
    def __init__(self, vocab_len, embedding_dim, hidden_dim, n_layers, dropout_prob):
        super().__init__()

        self.embedding = nn.Embedding(vocab_len, embedding_dim)
        self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout_prob)

        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, input_batch):
        embed = self.dropout(self.embedding(input_batch))
        outputs, (hidden, cell) = self.rnn(embed)

        return hidden, cell
```



# Many-to-Many: Machine translation example

```
class OneStepDecoder(nn.Module):
    def __init__(self, input_output_dim, embedding_dim, hidden_dim, n_layers, dropout_prob):
        super().__init__()
        # self.input_output_dim will be used later
        self.input_output_dim = input_output_dim

        self.embedding = nn.Embedding(input_output_dim, embedding_dim)
        self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout_prob)
        self.fc = nn.Linear(hidden_dim, input_output_dim)
        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, target_token, hidden, cell):
        target_token = target_token.unsqueeze(0)
        embedding_layer = self.dropout(self.embedding(target_token))
        output, (hidden, cell) = self.rnn(embedding_layer, (hidden, cell))

        linear = self.fc(output.squeeze(0))

        return linear, hidden, cell
```

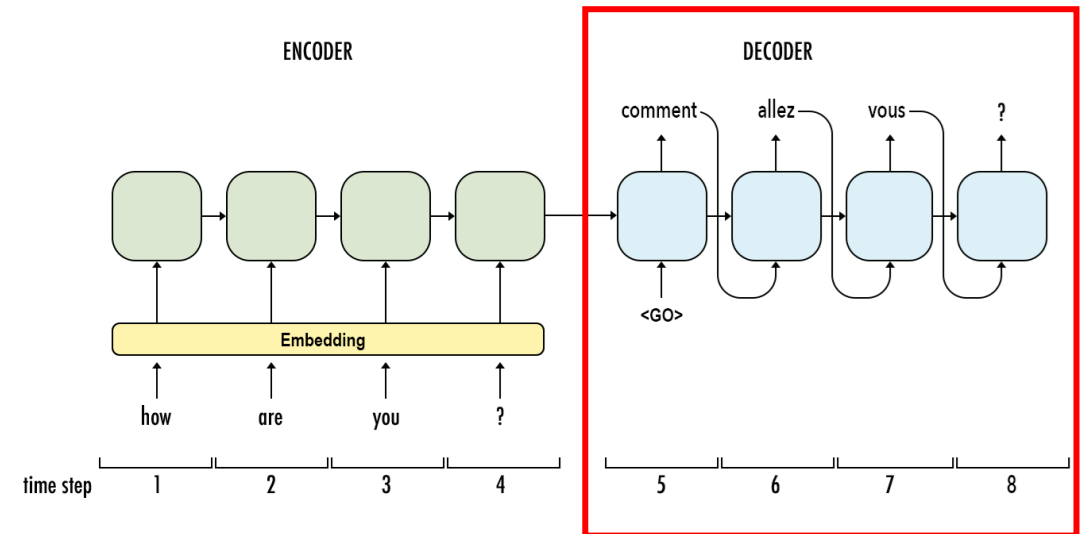
```
class Decoder(nn.Module):
    def __init__(self, one_step_decoder, device):
        super().__init__()
        self.one_step_decoder = one_step_decoder
        self.device=device

    def forward(self, target, hidden, cell):
        target_len, batch_size = target.shape[0], target.shape[1]
        target_vocab_size = self.one_step_decoder.input_output_dim
        # Store the predictions in an array for loss calculations
        predictions = torch.zeros(target_len, batch_size, target_vocab_size).to(self.device)
        # Take the very first word token, which will be sos
        input = target[0, :]

        # Loop through all the time steps
        for t in range(target_len):
            predict, hidden, cell = self.one_step_decoder(input, hidden, cell)

            predictions[t] = predict
            input = predict.argmax(1)

        return predictions
```




Implementation detail :

<http://www.adeveloperdiary.com/data-science/deep-learning/nlp/machine-translation-recurrent-neural-network-pytorch/>




# Part 3: Other RNN Variants


# Deep RNN and Bi-directional RNN



**Regular RNN**




**Deep RNN**



**Bi-directional RNN**

# DRNN: Pros and Cons



**Deep RNN**


(+)

Can provide better performance  
Often used for complex problems

(-)

Potential for overfitting  
Longer training time

# Bi-RNN: Pros and Cons



**Bi-directional RNN**

(+)

Higher performance in NLP tasks  
Suitable when both left and right contexts are used

(-)

Harder to train than Uni-directional RNN  
Not suitable for real-time processing

# DRNN/Bi-RNN Implementation


```
self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers = 2, dropout=dropout_prob, bidirectional = True)
```

Number of recurrent layers


Enable bi-directionality

# Lab Assignment: Predict stock prices using RNNs


# Dataset



**Tesla**



**Google**



**Dow Jones**

**Task:** Use many-to-one RNN architecture of choice to predict 100 days of stock values.

**Evaluation:** Plot the ground truth and predicted values for the last 100 days.